

OS – SEMINAR 1

1. Attendance requirements (see course website)

Commands and Paths

1. A command is a program, any program
2. We will only use commands that work in the console, meaning the interface is exclusively text
3. To run a command, type its name followed by whatever arguments necessary, everything separated by space. For instance:
 - a. To list the content of the current directory run `ls`
 - b. To see the content of file `/etc/passwd`, use `cat /etc/passwd`. Here, `/etc/passwd` is an argument
 - c. To see the content of the current directory, run `ls -l`. Here, `-l` is an argument too
 - d. To see the content of the current directory, with all the details, and including the hidden files, run `ls -l -a` or `ls -l --all`. Here, `-a` and `--all` have the same effect, one being the short form, and the other being the long form.
 - e. To create a directory name `abc`, run `mkdir abc`
 - f. To display the content of the current directory, with all the details but without the annoying colors, run `ls -l --color=never`. Here, `--color=never` is an argument with value. Sometimes, the equal sign is not necessary, but always consult the manual (command `man`) or the `--help` option (eg `ls --help`)
 - g. To do the same thing above, for the directory `/etc`, run `ls -l --color=never /etc`
4. Paths
 - a. UNIX file system has a single root, unlike Windows which has a root for every drive mounted (ie C:, D:, etc)
 - b. The UNIX file system root is `/`, and all drives are mounted as directories, somewhere in the file system
 - c. The UNIX file separator is `/`, unlike Windows, where the separator is `\`
 - d. Every user has a home directory, which is the current directory when you connect over SSH. Run command `pwd` to find the path to your current directory.

Seminar 1

1. Go through the ideas above explaining the command structure
 - a. Space is separator
 - b. First word is the command
 - c. Next words are arguments

- i. Values: ls /etc
- ii. Options
 - 1. Short form: ls -l
 - 2. Long form: ls --all
 - 3. Short form with value: cut -d : -f 1,2,3 /etc/passwd
 - 4. Long form with value: cut --delimiter=: --fields=1,2,3 /etc/passwd

Regular Expressions

- Very weird and ugly, but compact and flexible language for matching text.
- Why use them?
 - How can you find all the lines of a text file that contains phone numbers?
 - How can you remove all the extra spaces at the end of each line of a text file?
 - How can you remove any duplicated spaces from a text file?
 - If you ask a user to input an email in a text field, how do you verify that they input something that at least looks like an email?
- Extended Regular expression rules (use with option – E for grep and sed)
 - Every character that appears in a regular expression can have two meanings, normal or special, depending on the escape character \ appearing in front of it.
 - Depending on the program used to process the regular expressions, a characters special meaning is achieved with or without escaping it. Below are the meanings as required by the programs we will use for this class.

.	Matches any single character
\	Escape, changes the meaning of the character following it, between normal and special
[abc]	Matches any single character that appears in the list (in this case a or b or c)
[a-z]	Matches any single character that belongs to the range (in this case any lower-case letter)
[^0-9]	Matches any single character that does not belong to the range (in this case anything that is not a digit)
^	Beginning of line
\$	End of line

\<	Beginning of word
\>	End of word
()	Group several characters into an expression
*	Previous expression zero or more times
+	Previous expression one or more times
?	Previous expression zero or one times
{m,n}	Previous expression at least m and at most n times
	Logical OR between parts of the regular expression

Seminar 1

- Go through the aspects above, giving the following examples for the rules
 - .^{*} – any sequence of characters
 - [a-zA-Z02468] – any letter, regardless of its case, and any even digit
 - [!] – space or !
 - ^[^0-9]⁺\$ – non-empty lines containing any characters except digits
 - \([Nn][Oo]⁺ - any refusal, no matter how insistent (eg No no no no no)

Grep

- Program that searches through files using regular expressions
- Why is it called `grep` and not something more intuitive? It was very intuitive to those who implemented `grep`, but not so much to us. Google “why is grep called grep” and enjoy. While at it, go learn about `ed` (the original UNIX editor). It is available in your Linux installation. Have Google handy if you play with it ...
- Option arguments that we will use a lot
 - `-v` – display lines that do not match the given regular expression
 - `-i` – ignore upper/lower case when matching (match case-insensitively)
 - `-q` – Do not display the matching lines, just exit with 0 if found, or 1 if not found. We will use this only later , when we get to Shell Programming.
 - `-E` to use extended regular expressions
 - `-c` – count lines matching
 - `-o` – only matching part of a line, each on a separate line

4. File `/etc/passwd` contains all the users in the system, providing their usernames, full names, home directories, etc
- We will use this file a lot to exercise text matching
 - It is structured as follows, using `:` as field separator
`username:password:user-id:group-id:user-info:home-directory:shell`
 - The password field will always be `x`, the actual encrypted password residing in file `/etc/shadow` which cannot be read by regular users
 - The user-info field, usually contains the full name of the user

Seminar 1

- Let's search for things in file `/etc/passwd`
 - Display all lines containing "dan". The solution is below
 - `grep "dan" /etc/passwd`
 - Display the line of username "dan". The username is the first field on the line, it is not empty, and it ends at the first `:`. We will rely on these aspects to insure that we only search the usernames, and not anything else.
 - `grep -i "^dan:" /etc/passwd`
 - Display the lines of all users who do not have digits in their username.
 - `grep -E "^[^0-9:]+:" /etc/passwd`
 - Display the lines of all users who have at least two vowels in their username. This is a little tricky, because the vowels do not need to be consecutive, so we need to allow for any characters between the vowels (including none), but we cannot allow `:` to appear between vowels, or else we would be searching outside the username.
 - `grep -i -E "^[^:]*[aeiou][^:]*[aeiou][^:]*:" /etc/passwd`
 - `grep -i -E "^[^:]*([aeiou][^:]*){2,}:" /etc/passwd`
 - There will be lots of users displayed for the problem above, so let's search for usernames with at least 5 vowels in their username. The first solution above will be really long for this case, but the second will be very easy to adapt, by changing 2 into 5.
 - `grep -E -i "^[^:]*([aeiou][^:]*){5,}:" /etc/passwd`
 - Display the lines of all the users not having bash as their shell. The shell is the last value on the line, so we will use that when searching.
 - `grep -v "/bash$" /etc/passwd`
 - Display the lines of all users named Ion. We will have to search in the user-info field (the fifth field) of each line, ignore the upper/lower case of the letters, and insure that we do not display anybody containing the sequence "ion" in their names (eg Simion, Simionescu, or Ionescu).
 - `grep -E -i "^[^:]*(:){4}[^:]*<ion>" /etc/passwd`
- Let's consider a random text file `a.txt`, and search for things in it
 - Display all the non-empty lines
 - `grep "." a.txt`
 - Display all the empty lines

- i. `grep "^$" a.txt`
- c. Display all lines containing an odd number of characters
 - i. `grep -E "^(..)*.$" a.txt`
- d. Display all lines containing an ocean name
 - i. `grep -E -i "\<atlantic\>|\<pacific\>|\<indian\>|\<arctic\>|\<antarctic\>" a.txt`
- e. Display all lines containing an email address
 - i. What does an email address look like? It has the following structure.
 - 1. username – let's assume it can contain any character, except for @, *, !, and ?
 - 2. @ - separator between the username and the hostname
 - 3. hostname
 - a. Sequence of at least two elements separated by .
 - b. Let's assume an element can contain any letter, digit, dash, or underscore
 - ii. `grep -E -i "\<[^@*\\!?!]+@[a-z0-9_-]+(\.[a-z0-9_-]+)+\>" a.txt`

Sed

1. Program for searching processing text by performing search/replace, transliterations, line deletion, etc
2. By default, it does not modify the file, but displays the result of processing the input file
3. We will use only the features presented below
 - a. Search/replace
 - i. Command structure: `sed "s/regex/replacement/flags" a.txt`
 - ii. `s` – is the search/replace command
 - iii. `/` is the separator, and can be any other character. The first character after the command `s` is considered to be the separator
 - iv. The flags at the end can be `g`, `i`, or both
 1. `g` – Perform the replacement everywhere on the line. Without it, only the first appearance will be replaced
 2. `i` – Perform a case-insensitive search
 - v. The replacement can contain reference to the expressions grouped in the regex as `\1`, `\2`, etc, the number being the order in which the groups appear in the regex
 - b. Transliterate
 - i. Command structure: `sed "y/characters/replacement/" a.txt`
 - ii. `y` – is the transliteration command
 - iii. `/` is the separator, and can be any other character. The first character after the command `y` is considered to be the separator

iv. The characters and the replacement must have the same length

c. Delete lines matching a regular expression

- i. Command structure: `sed "/regex/d" a.txt`
- ii. / is the separator
- iii. d – is the line deletion command

Seminar 1

1. Let's manipulate the content of /etc/passwd

a. Display all lines, replacing all vowels with spaces

- i. `sed "s/[aeiou]/ /gi" /etc/passwd`

b. Display all lines, converting all vowels to upper case

- i. `sed "y/aeiou/AEIOU/" /etc/passwd`

c. Display all lines, deleting those containing numbers of five or more digits:

- i. `sed -E "[0-9]{5,}/d" /etc/passwd`

d. Display all lines, swapping all pairs of letters

- i. `sed -E "s/([a-z])([a-z])/\2\1/gi" /etc/passwd`

e. Display all lines, duplicating all vowels

- i. `sed -E "s/([aeiou])/\1\1/gi" /etc/passwd`

Awk

1. Given a separator character (by default it is space), treats the input text as a table, with each line being a row, and the fields of each row the tokens of the line, as determined by the separator.
2. Processes the input based on a program written in a simple C-like language
3. A program is a sequence of instruction blocks, prefixed by an optional selector
4. Each block in the program is applied to every line of input matching its selector. If the block does not have a selector, it is applied to every line of input
5. A selector is any valid conditional expression, or one of the following two special selectors
 - a. BEGIN – the block associated with this selector is executed before any input has been processed
 - b. END – the block associated with this selector is executed after all input has been processed
6. Special variables
 - a. NR – number of the current line of input
 - b. NF – the number of fields on the current line
 - c. \$0 – the entire input line
 - d. \$1, \$2, ... – the fields of the current line

7. The AWK program can be written in a file, or provided directly on the command line between apostrophes

1.1.1 Seminar 1

1. Manipulate the content of `/etc/passwd`, using AWK with the program provided on the command line
 - a. Display all the usernames, but only the usernames, and nothing else. We will use argument `-F` to tell AWK that the input file is separated by `:`, and then we will print the first field of each line, by not providing any selector for the block.
 - i. `awk -F: '{print $1}' /etc/passwd`
 - b. Print the full name (the user info field) of the users on odd lines
 - i. `awk -F: 'NR % 2 == 1 {print $5}' /etc/passwd`
 - c. Print the home directory of users having their usernames start with a vowel
 - i. `awk -F: '/^[aeiouAEIOU]/ {print $6}' /etc/passwd`
 - d. Print the full name of users having even user ids
 - i. `awk -F: '$3 % 2 == 0 {print $5}' /etc/passwd`
 - e. Display the username of all users having their last field end with "nologin"
 - i. `awk -F: '$NF ~ /nologin$/ {print $1}' /etc/passwd`
 - f. Display the full names of all users having their username longer than 10 characters
 - i. `awk -F: 'length($1) > 10 {print $5}' /etc/passwd`
2. Keep using `/etc/passwd` as input file, but provide AWK programs in a file. The command will look like
 - a. `awk -F: -f prog.awk /etc/passwd`
 - b. Provide the content of file `prog.awk` so that the command above will print all user on even line having a group id less than 20

```
NR % 2 == 0 && $4 < 20 {  
    print $5  
}
```

- c. Display the sum of all user ids

```
BEGIN {  
    sum=0  
}  
{  
    sum += $3  
}  
END {  
    print sum  
}
```

- d. Display the product of the differences between the user id and the group id

```
BEGIN {  
    prod=1  
}  
{  
    prod *= $3-$4  
}  
END {  
    print prod  
}
```