

Local Network - collection of computers interconnected

- components:
  - hosts (computers)
  - links
  - switches / routers / hubs
  - access points

unit of communication = frame (electrical signal)

LAN = only switches & hubs

switch - amount of bit - chat is minimum

hub - signal is sent to everyone; scrambled signals can happen, collision spam is larger

switches are usually better

- look up MAC addresses (physical) from the factory

host puts in the frame source and destination MAC address

LEDs signaling collisions  $\Rightarrow$  resend the signal

broadcast: computer talking to all computers linked to a switch (hub-like behaviour)

unicast: two hosts talking to each other

destination MAC for broadcast: FF FF FF FF FF FF

hub bandwidth < switch bandwidth

multicast: one host talking to a controlled subset of machines

broadcast is only possible in LAN, while multicast is possible out of LAN

routers - interconnect networks

frame - contains IP addresses in the datagram

## TCP/IP

- host needs identifier  $\Rightarrow$  unique IP address

- port - identifies process; ex: HTTP 80, HTTPS 443

data loss, timing (delay), bandwidth

TCP - reliable transport, no data is lost, congestion control, no timing

UDP - unreliable transport, no timing, possible data loss, but very quick

TCP regulates traffic, UDP does not

there are apps that tolerate not having those things, like Skype's call part

socket - OS controlled interface

TCP:

- server listening for clients and must have a socket welcoming the client
- client creates TCP socket and binds to it on IP and a port and connects to it
- creating socket is immediate, while accept, receive are blocking calls

1-1024 ports are reserved

AF\_INET - address family (TCP/IP)

INADDR\_ANY = 0.0.0.0 IP address is 4 bytes  $\Rightarrow$  integer

listen = creates a queue; at most that number of connections

loopback - allows to run network apps while not connected

send, recv - TCP/IP specific and give greater control to the programmer

write, read - more general

netstat to see used ports

socket types: UDP = SOCK\_DGRAM, TCP = SOCK\_STREAM

sockaddr has the field sa\_data[14], which is opaque

this gets split into 3 parts in sockaddr\_in: sin\_port (2 bytes), sin\_addr (4 bytes) and sin\_zero [8] not used

union in C: chunk of memory used to store several variables (C polymorphism)

in\_addr is a union; it can be filled with anything we want on 4 bytes

send returns the number of bytes sent and successfully received, recv returns the number of bytes received

both return -1 in case of an error

shutdown - close one of the send / receive ends of a socket



## TCP - connection oriented

- whatever you send is received, like a queue

- guaranteed data delivery + guaranteed data ordering delivery

server types - iterative servers (blocking) - handle clients one after another

- concurrent servers (blocking) - fork, threads

- concurrent multiplexed servers (select)

blocked - until a call fails or succeeds, the app is stuck at that part of the code

dns: gethostname, gethostbyaddr return hostname

hostname has a list of addresses and a list of aliases

a machine may have multiple IP addresses

uint 32 - to ensure integers are 4 bytes

big endian = most significant byte first little endian least significant byte first

always send on socket on big endian because you don't know the architecture of the other machine

float to network: look at it like it's an integer and apply htonl

double to network: swap the bytes manually by looking at it like a char array

## UDP - connectionless, datagram oriented

datagram = basic transfer unit; individual messages

recvfrom, sendto must specify the destination address

read can be used as well

no guarantee for delivery or ordering

less overhead than TCP, less latency, higher bandwidth, can broadcast

packet = more 64kb, send more  $\Rightarrow$  the rest is lost

CRC = cyclic redundancy check; it is a checksum

one party can overflow the other

no listen, no accept because there's no connection

sendto returns -1 only in case of locally detected errors

recvfrom matches 1 recvfrom

each packet has a priority; routers are the ones discarding packets

recvfrom is blocked until you get some data

socket options - SO\_REUSEADDR - reuse local addresses; get rid of addresses after interrupting

SO\_BROADCAST - enables broadcast

blocking  $\Rightarrow$  process is put to sleep

not socket to nonblocking  $\Rightarrow$  calls won't block, but will return EWOULDBLOCK instead

multiplexed (select) - allows to watch over multiple file descriptors, and be notified when something

changes; process is sleeping for a controlled interval of time

signal i/o model - signaled when data is available

async read - returns control but data is being waited for

select - watches sockets for read/write/exceptions for a given interval of time

- returns -1 on error, 0 on timeout and positive count of ready descriptors otherwise

- select is destructive with respect to the given sets

- NULL timeout  $\Rightarrow$  wait indefinitely

broadcast does not work on TCP because you must have an existing connection

broadcast destination address [255.255.255.255 (universal broadcast)

192.168.0.255 (LAN broadcast) = largest IP

nc - send pings

there can be an unlimited number of broadcasters

ping - Windows sends 4 pings, Linux sends until it is stopped

- checks whether a machine is up and running; you can ping a broadcast address

netstat - sockets connected

tracert - route packets



knowledge flag  
reset flag = 1 → connect to TCP  
protocol - an agreement about communication  
- specifies format, meaning, rules of exchange and error handling  
- possible problems: corrupted bits, entire packets lost, duplicated packets, packets out of order

network = stack of layers  
architecture = set of layers and protocols  
layers: physical, data link, network, transport, session, presentation, application

physical layer: transmits a stream of raw bits over a communication channel  
unit of transmission: signal  
data link layer: turns raw transmissions into error-free communication  
handles flow control

network layer: unit of transmission: frames = thousand of bytes  
controls the operation of a subnet  
routes packets from a source to a destination  
handles congestion control (saturated router)  
quality of service = data priorities  
fragmentation = when data to be sent is too big

transport layer: accepts data from upper layers and splits it into packets  
inserts that packets arrive correctly → reliability  
end-to-end  
unit of transmission: packets (UDP), segments / sessions (TCP)

session layer: allows establishing sessions  
unit of transmission: control message

presentation layer: - data translator for the network  
- ntohs, htons, string encoding happens there  
application layer: - this is what we implement  
- unit of transmission: message (user data)

OPEN SYSTEM INTERCONNECTION  
TCP/IP has 4 layers → mappings are not 1 to 1  
RFC - request for comments

Network layer - Internet Protocol (IP)

IP address: 32-bit identifier for host/router interface →  $2^{32}$  IP addresses  
interface: connection between host/router and physical link

LAN = can physically reach each other without a router  
all routers have a routing table → destination, gateway, interface

if the IP addresses were unstructured → routing tables  $\approx 2^{32}$  entries → it can be very slow  
5 classes of networks:

A: 1 byte network, 3 bytes hosts, network prefix: 0

B: 2 bytes network, 2 bytes hosts, network prefix: 10

C: 3 bytes network, 1 byte hosts, network prefix: 110

D: multicast with prefix 1110

E: experimental with prefix 1111

255.255.255.255 is not used

classful - allocate to a network an entire class

network address - smallest IP in the network

broadcast address - greatest IP in the network

classless interdomain routing

address format: a.b.c.d/x, x is the number of network bits

reduces sizes of routing tables because of subnetting

the beginning address must be divisible by the number of addresses in the block

masks are flexible → networks can be divided into smaller subnets

all classes have natural masks

classful is limited

we cannot have subnets with 2 IPs

supernetting: multiple routing entries become one



127.0.0.0/8 - local host  
0.0.0.0 is the default route for a router  
IP datagram: fields, source and destination IP addresses, data  
data link: MAC addresses  
network: IP addresses  
transport: ports

### ARP - address resolution protocol

- when sending a packet to another network, IP destination is the same, but the MAC address is that of the router

### IP datagram

- 20 bytes of IP + data
  - version (4 bits) + header length (4 bits) + type of service (1 byte) + length (2 bytes)
  - 16 bit identifier (2 bytes) + 13 bit offset for fragments + DF/MF (don't fragment) more fragments
  - TTL (time to live, 1 byte) - to avoid packets that never die, decremented on each router
  - upper layer protocol (1 byte) - UDP, TCP, ICMP
  - checksum (2 bytes) - to check if there are any errors
  - each router recomputes the checksum because of TTL
  - checksum = 16 bit 1's complement of the 1's complement sum of all 16-bit words in the header
  - fragmentation: networks have MTU (maximum transfer unit)
  - regular ethernet: 1500 → 9000 bytes (jumbo frames)
  - IP splits data when necessary, but reconstructs it only at the destination
  - offset: MTU - 20 bytes
- ARP broadcasts with source IP address, destination IP address, its MAC address and an empty MAC address and waits for it to be filled

### NAT - network address translation

addresses are hidden

- private IPs - no router in internet will route them
- go out in the internet with a private IP
- router changes packet IP address to its public address

### NAT translation table

- a different port for each different connection
- one IP address for all services ⇒ local addresses can change
- devices inside network not explicitly addressable
- ~60000 simultaneous connections for UDP + another 60000 for TCP
- routers should only process up to layer 3
- address shortage should be solved by IPv6

### ICMP - signal things, error reporting (for UDP/TCP), echo request

- 4 bytes header: 1 byte type, 1 byte code, 2 bytes checksum

### traceroute

- UDP misuse, ICMP windows
- uses IP TTL to trace paths ⇒ TTL=1 get first router; TTL=2 get second router; ...
- uses highly unlikely to be used UDP port

### UDP

- source port and destination port (2 bytes each)
- length (2 bytes) ≥ 8, checksum for the entire datagram (2 bytes)
- 28 bytes + data

### TCP

- each segment needs confirmation, info network and peer status
- source port and destination port (2 bytes each)
- sequence number (4 bytes) - incremented for each segment sent
  - allows segments to be sent in order
- acknowledgement number (4 bytes) - helps confirm all segments we sent
- TCP header length + not used area + flags (2 bytes)



network = stack of protocols, format, meaning, ...  
lower archite.

- acknowledge flag
- reset flag = 1 → connect to TCP & no one listening
- syn, fin = 1 when setting and shutting down a connection
- window size (2 bytes) - how much data you're willing to accept
- urgent pointer (2 bytes) - not used
- options - setsockopt

### TCP segments:

- data is split into segments; each segment is carried into one or multiple IP datagrams
- each needs acknowledgment
- buffers for sending and receiving
- state: seq. no. - amount of bytes sent  
last acknowledged byte

- connect: request, granted, ack  
seq - random number (local 0)  
syn = 1 during connection establishment  
ack = 1 acknowledge I got the state variables

- teardown: 4 steps  
client closes its half (FIN), server acknowledges it  
server sends the remaining data, server closes its half (FIN), client acknowledges it

- sequence numbers:
  - used to reassemble data
  - increments based on the number of bytes in the TCP data field
  - ACK indicates the next byte number to receive

- sliding window:
  - all data in the window is allowed to be sent without an acknowledgement
  - when acknowledgement arrives, the window is shifted
  - if we sent all data and no acknowledgement has been received, TCP stops

### receiver buffer:

- app may not be reading data as it comes
- window size: still free space
- sender is allowed to send the minimum between its sliding window size and the advertised window size
- sender keeps sending TCP segments of size  $s$  when the advertised window is 0 (flow control)

TCP will retransmit a segment upon expiration of a timer  
round trip time - send segment and await acknowledgement  
estimated RTT =  $(1-\alpha)$  estimated RTT +  $\alpha$  sample RTT  $\alpha \in (0,1), \alpha = 0.125$   
dev RTT =  $(1-\beta)$  dev RTT +  $\beta$  sample RTT - estimated RTT  
timeout interval = estimated RTT + 4 dev RTT

first acknowledgement is lost  $\Rightarrow$  it is fine as long as the second one is received  
congestion: too many sources sending too much data  
- lost packets, long delays, router just starts dumping packets

end-to-end congestion - congestion is inferred

### congestion window:

- starts with one segment
- TCP can send min (sliding window, advertised window, congestion window)
- congestion window is doubled until first loss
- loss: timeout or 3 duplicate acks
- ack is cumulative, but doesn't allow gaps
- 3 dup. acks: window is halved, then it grows linearly (I have some loss, but data is still travelling)
- timeout: window is set to 1 (this is bad)
- TCP is fair: all computers use the same bandwidth approximately



## routing

- determine good paths
- static or dynamic
  - static: routes change slowly
  - dynamic: routes change quickly, periodic updates
- global: all routers have complete info
- decentralized: router only knows its neighbours

## Routing Information Protocol (RIP)

- each router sends the routing table to its neighbours regularly and updates it accordingly
- v1 - classful addressing
- v2 - CIDR
- when a link is down, there may be a loop happening = counting to infinity
- RIP v1 - max 15 to limit counting to infinity
- RIP v2 - poisoned reverse - A does not send entries to B passing through B

## Application Level Protocols

- DNS - domain name server
- names are easier to remember
- distributed database of domain addresses and their IP addresses
- hierarchical naming system
- host.name.domain.TLD
  - TLD - top level domain standardised
  - domain + TLD - name that you buy

## DNS Software

- resolver: name  $\rightarrow$  info
- name server: DNS database
- originally one single central table

## Servers

- primary / master - authoritative, updates regularly
- secondary / slaves - temporarily authoritative, read-only copies of the database
- forwarders - resolver only
- queries - recursive or non-recursive

## types of records:

- A - host to IP
- NS - who is the name server for a domain
- CNAME - alias
- MX - mail

TTL - time a record should be kept alive

DNS - 12 bytes UDP

## FTP Protocol

- clear text, text protocol, exchanging files between machines
- 2 channels: control channel + data channel
- modes: active, passive
- TCP

## SMTP Protocol

- mail exchange linked to DNS
- text protocol which allows encryption (TCP) and offline message exchanging

## HTTP Protocol

- text protocol, TCP, which allows you to connect to a website