

Databases

Lecture 7

Relational Algebra

- query languages in the relational model
 - relational algebra and calculus - formal query languages with a significant influence on SQL
 - relational algebra
 - queries are specified in an operational manner
 - relational calculus
 - queries describe the desired answer, without specifying how it will be computed (declarative)
 - not expected to be Turing complete
 - not intended for complex calculations
 - provide efficient access to large datasets
 - allow optimizations

- relational algebra
 - used by DBMSs to represent query execution plans
 - a relational algebra query:
 - is built using a collection of operators
 - describes a step-by-step procedure for computing the result set
 - is evaluated on the input relations' instances
 - produces an instance of the output relation
 - every operation returns a relation, so operators can be composed; the algebra is closed
 - the result of an algebra expression is a relation, and a relation is a set of tuples
- relational algebra on bags (multisets) - duplicates are not eliminated

Conditions

- conditions that can be used in several algebraic operators
- similar to the SELECT filter conditions

1. *attribute_name relational_operator value*

- *value* - attribute name, expression

2. *attribute_name IS [NOT] IN single_column_relation*

- a relation with one column can be considered a set
- the condition tests whether a value belongs to a set

3. *relation {IS [NOT] IN | = | <>} relation*

- the relations in the condition must be union-compatible

Conditions

4. *(condition)*

NOT condition

condition₁ AND condition₂

condition₁ OR condition₂,

where *condition*, *condition₁*, *condition₂* are conditions of type 1-4.

Operators in the Algebra

- equivalent SELECT statements can be specified for the relational algebra expressions
- *selection*
 - notation: $\sigma_C(R)$
 - resulting relation:
 - schema: R 's schema
 - tuples: records in R that satisfy condition C
 - equivalent SELECT statement

```
SELECT *  
FROM R  
WHERE C
```

- *projection*
 - notation: $\pi_{\alpha}(R)$
 - resulting relation:
 - schema: attributes in α
 - tuples: every record in R is projected on α
 - α can be extended to a set of expressions, specifying the columns of the relation being computed
 - equivalent SELECT statement

```
SELECT DISTINCT  $\alpha$ 
FROM R
```

```
SELECT  $\alpha$ 
FROM R                -- algebra on bags
```

- *cross-product*
 - notation: $R_1 \times R_2$
 - resulting relation:
 - schema: the attributes of R_1 followed by the attributes of R_2
 - tuples: every tuple r_1 in R_1 is concatenated with every tuple r_2 in R_2
- equivalent SELECT statement

```
SELECT *  
FROM R1 CROSS JOIN R2
```


- *union, set-difference, intersection*
 - notation: $R_1 \cup R_2$, $R_1 - R_2$, $R_1 \cap R_2$
 - R_1 and R_2 must be union-compatible:
 - same number of columns
 - corresponding columns, taken in order from left to right, have the same domains
 - equivalent SELECT statements

SELECT *	SELECT *	SELECT *
FROM R1	FROM R1	FROM R1
UNION	EXCEPT	INTERSECT
SELECT *	SELECT *	SELECT *
FROM R2	FROM R2	FROM R2

-- algebra on bags: SELECT statements that don't eliminate duplicates (e.g., UNION ALL)

- join operators
 - *condition join* (or *theta join*)
 - notation: $R_1 \otimes_{\Theta} R_2$
 - result: the records in the cross-product of R_1 and R_2 that satisfy a certain condition
 - definition $\Rightarrow R_1 \otimes_{\Theta} R_2 = \sigma_{\Theta}(R_1 \times R_2)$
 - equivalent SELECT statement


```
SELECT *
FROM R1 INNER JOIN R2 ON  $\Theta$ 
```

- join operators
 - *natural join*
 - notation: $R_1 * R_2$
 - resulting relation:
 - schema: the union of the attributes of the two relations (attributes with the same name in R_1 and R_2 appear once in the result)
 - tuples: obtained from tuples $\langle r_1, r_2 \rangle$, where r_1 in R_1 , r_2 in R_2 , and r_1 and r_2 agree on the common attributes of R_1 and R_2
 - let $R_1[\alpha]$, $R_2[\beta]$, $\alpha \cap \beta = \{A_1, A_2, \dots, A_m\}$; then:

$$R_1 * R_2 = \pi_{\alpha \cup \beta} (R_1 \otimes_{R_1.A_1=R_2.A_1 \text{ AND } \dots \text{ AND } R_1.A_m=R_2.A_m} R_2)$$
 - equivalent SELECT statement


```
SELECT *
FROM R1 NATURAL JOIN R2
```

- join operators
 - *left outer join*
 - notation (in these notes): $R_1 \bowtie_C R_2$
 - resulting relation:
 - schema: the attributes of R_1 followed by the attributes of R_2
 - tuples: tuples from the condition join $R_1 \bowtie_C R_2$ + the tuples in R_1 that were not used in $R_1 \bowtie_C R_2$ combined with the *null* value for the attributes of R_2
 - equivalent SELECT statement

```
SELECT *  
FROM R1 LEFT OUTER JOIN R2 ON C
```

- join operators
 - *right outer join*
 - notation: $R_1 \bowtie_C R_2$
 - resulting relation:
 - schema: the attributes of R_1 followed by the attributes of R_2
 - tuples: tuples from the condition join $R_1 \Join_C R_2$ + the tuples in R_2 that were not used in $R_1 \Join_C R_2$ combined with the *null* value for the attributes of R_1
 - equivalent SELECT statement

```
SELECT *  
FROM R1 RIGHT OUTER JOIN R2 ON C
```

- join operators
 - *full outer join*
 - notation: $R_1 \bowtie_C R_2$
 - resulting relation:
 - schema: the attributes of R_1 followed by the attributes of R_2
 - tuples:
 - tuples from the condition join $R_1 \otimes_C R_2$ +
 - the tuples in R_1 that were not used in $R_1 \otimes_C R_2$ combined with the *null* value for the attributes of R_2 +
 - the tuples in R_2 that were not used in $R_1 \otimes_C R_2$ combined with the *null* value for the attributes of R_1
 - equivalent SELECT statement


```
SELECT *
FROM R1 FULL OUTER JOIN R2 ON C
```

- join operators
 - *left semi join*
 - notation: $R_1 \triangleright R_2$
 - resulting relation:
 - schema: R_1 's schema
 - tuples: the tuples in R_1 that are used in the natural join $R_1 * R_2$

- join operators
 - *right semi join*
 - notation: $R_1 \Join R_2$
 - resulting relation:
 - schema: R_2 's schema
 - tuples: the tuples in R_2 that are used in the natural join $R_1 * R_2$

- *division*

- notation: $R_1 \div R_2$
- $R_1[\alpha], R_2[\beta], \beta \subset \alpha$
- resulting relation:
 - schema: $\alpha - \beta$
 - tuples: a record $r \in R_1 \div R_2$ iff $\forall r_2 \in R_2, \exists r_1 \in R_1$ such that:
 - $\pi_{\alpha-\beta}(r_1) = r$
 - $\pi_{\beta}(r_1) = r_2$
 - i.e., a record r belongs to the result if in R_1 r is concatenated with every record in R_2



- see lecture examples (at the board) with algebra queries:
 - selection
 - projection
 - division
 - selection, projection
 - natural join, selection, projection
 - different algebra expressions producing the same result (optimization - reducing the size of intermediate relations)

An Independent Subset of Operators

- independent set of operators M:
 - eliminating any operator op from M: there will be a relation that can be obtained using M's operators, but cannot be obtained with the operators in $M - \{op\}$
- for the previously described query language, with operators:
 $\{\sigma, \pi, \times, \cup, -, \cap, \otimes, *, \ltimes, \rtimes, \bowtie, \triangleright, \triangleleft, \div\}$
an independent set of operators is $\{\sigma, \pi, \times, \cup, -\}$
- the other operators are obtained as follows (some expressions have already been introduced):
 - $R_1 \cap R_2 = R_1 - (R_1 - R_2)$
 - $R_1 \otimes_C R_2 = \sigma_C(R_1 \times R_2)$

- the other operators are obtained as follows (some expressions have already been introduced):

- $R_1[\alpha], R_2[\beta], \alpha \cap \beta = \{A_1, A_2, \dots, A_m\}$, then:

$$R_1 * R_2 = \pi_{\alpha \cup \beta}(R_1 \otimes_{R_1.A_1=R_2.A_1 \text{ AND } \dots \text{ AND } R_1.A_m=R_2.A_m} R_2)$$

- $R_1[\alpha], R_2[\beta], R_3[\beta] = \{(null, \dots, null)\}, R_4[\alpha] = \{(null, \dots, null)\}$

$$R_1 \bowtie_C R_2 = (R_1 \otimes_C R_2) \cup (R_1 - \pi_\alpha(R_1 \otimes_C R_2)) \times R_3$$

$$R_1 \bowtie_C R_2 = (R_1 \otimes_C R_2) \cup R_4 \times (R_2 - \pi_\beta(R_1 \otimes_C R_2))$$

$$R_1 \bowtie_C R_2 = (R_1 \bowtie_C R_2) \cup (R_1 \bowtie_C R_2)$$


- $R_1[\alpha], R_2[\beta]$

$$R_1 \triangleright R_2 = \pi_\alpha(R_1 * R_2)$$

$$R_1 \triangleleft R_2 = \pi_\beta(R_1 * R_2)$$

- the other operators are obtained as follows (some expressions have already been introduced):
 - if $R_1[\alpha]$, $R_2[\beta]$, $\beta \subset \alpha$, then $r \in R_1 \div R_2$ iff $\forall r_2 \in R_2, \exists r_1 \in R_1$ such that:
 $\pi_{\alpha-\beta}(r_1) = r$ and $\pi_\beta(r_1) = r_2$
 $\Rightarrow r$ is in $\pi_{\alpha-\beta}(R_1)$, but not all the elements in $\pi_{\alpha-\beta}(R_1)$ are in the result
 - $(\pi_{\alpha-\beta}(R_1)) \times R_2$ contains all the elements with one part in $\pi_{\alpha-\beta}(R_1)$ and the second part in R_2
 - to obtain values that are disqualified, R_1 is subtracted from the obtained relation, and the result is projected on $\alpha - \beta$
 - the final expression:

$$R_1 \div R_2 = \pi_{\alpha-\beta}(R_1) - \pi_{\alpha-\beta}((\pi_{\alpha-\beta}(R_1)) \times R_2 - R_1)$$



- the *renaming* operator

$$\rho(R'(A_1 \rightarrow A_1', A_2 \rightarrow A_2', A_3 \rightarrow A_3'), E)$$

- E - relational algebra expression
- the result, relation R', has the same tuples as the result of E
- attributes A₁, A₂, and A₃ are renamed to A₁', A₂', and A₃', respectively

* the next examples use the statements below:

- assignment

$R[\text{list}] := \text{expression}$

- the expression's result (a relation) is assigned to a variable ($R[\text{list}]$), specifying the name of the relation [and the names of its columns]

- eliminating duplicates from a relation

$\delta(R)$

- sorting records in a relation

$S_{\{\text{list}\}}(R)$

- grouping

$Y_{\{\text{list1}\} \text{ group by } \{\text{list2}\}}(R)$

- R 's records are grouped by the columns in *list2*
- *list1* (that can contain aggregate functions) is evaluated for each group of records

students [id, name, sgroup, gpa, dob]

groups [id, year, program]

schedule [day, starthour, endhour, activtype, room, sgroup,
facultym_id]

faculty_members [id, name]

1. The names of students in a given group:

$$R := \pi_{\{name\}} \left(\sigma_{sgroup='222'}(students) \right)$$

SELECT name

FROM students

WHERE sgroup='222'

2. The students in a given program (alphabetical list, by groups):

$$G := \pi_{\{id\}} \left(\sigma_{program='IG'}(groups) \right)$$
$$R := S_{\{sgroup, name\}} \left(\sigma_{sgroup \text{ is in } G}(students) \right)$$

```
SELECT *  
FROM students  
WHERE sgroup IN  
    (SELECT id  
     FROM groups  
     WHERE program='IG')  
ORDER BY sgroup, name
```

```
students [id, name, sgroup, gpa, dob]  
groups [id, year, program]  
schedule [day, starthour, endhour, activtype, room,  
          sgroup, facultym_id]  
faculty_members [id, name]
```

3. The number of students in every group of a given program:

$$ST := \sigma_{sgroup \text{ is in } \left(\pi_{\{id\}} \left(\sigma_{program='IG'}(groups) \right) \right)}(students)$$

$$NR := \gamma_{\{sgroup, count(*)\}} \text{ group by } \{sgroup\} (ST)$$

```
SELECT sgroup, COUNT(*)
FROM (SELECT *
      FROM students
      WHERE sgroup IN
        (SELECT id
         FROM groups
         WHERE program='IG')
      ) t
GROUP BY sgroup
```

students [id, name, sgroup, gpa, dob]
groups [id, year, program]
schedule [day, starthour, endhour, activtype, room,
sgroup, facultym_id]
faculty_members [id, name]

4. A student's schedule (the student is given by name):

$$T := \sigma_{sgroup \text{ is in } \left(\pi_{\{sgroup\}} \left(\sigma_{name='Ionescu M. Razvan'}(students) \right) \right)}(schedule)$$

5. The number of hours per week for every group:

$$F(no, sgroup) := \pi_{\{endhour - starthour, sgroup\}}(schedule)$$
$$NoHours(sgroup, nohours) := \gamma_{\{sgroup, sum(no)\}} \text{ group by } \{sgroup\}(F)$$

students [id, name, sgroup, gpa, dob]

groups [id, year, program]

schedule [day, starthour, endhour, activtype, room, sgroup, facultym_id]

faculty_members [id, name]

6. The faculty members (their names) who teach a given student:

$$A := (\sigma_{name='Ionescu M. Razvan'}(students)) \otimes_{students.sgroup=schedule.sgroup} schedule$$
$$B := \pi_{\{faculty_id\}}(A)$$
$$C := faculty_members \otimes_{faculty_members.id=B.facultym_id} B$$
$$D := \pi_{\{name\}}(C)$$

students [id, name, sgroup, gpa, dob]

groups [id, year, program]

schedule [day, starthour, endhour, activtype, room, sgroup, facultym_id]

faculty_members [id, name]

7. The faculty members with no teaching assignments (i.e., not on the schedule):

$$C := \pi_{\{name\}}(faculty_members) - \pi_{\{name\}}(schedule \otimes_{schedule.facultym_id=faculty_members.id} faculty_members)$$

* Is there a problem if two different faculty members have the same name?

students [id, name, sgroup, gpa, dob]

groups [id, year, program]

schedule [day, starthour, endhour, activtype, room, sgroup, facultym_id]

faculty_members [id, name]

8. Students with school activities on every day of the week (all days with school activities considered):

$$A := \delta \left(\pi_{\{day\}}(schedule) \right)$$

$$B := students \otimes_{students.sgroup=schedule.sgroup} schedule$$

$$C := \delta \left(\pi_{\{name, day\}}(B) \right)$$

$$D := C \div A$$

* Is there a problem if two different students have the same name?

students [id, name, sgroup, gpa, dob]

groups [id, year, program]

schedule [day, starthour, endhour, activtype, room, sgroup, facultym_id]

faculty_members [id, name]

References

- [Ta13] ȚÂMBULEA, L., Curs Baze de date, Facultatea de Matematică și Informatică, UBB, 2013-2014
- [Ra02] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (3rd Edition), McGraw-Hill, 2002
- [Da03] DATE, C.J., An Introduction to Database Systems (8th Edition), Addison-Wesley, 2003
- [Ga09] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book (2nd Edition), Pearson Education, 2009
- [Ha96] HANSEN, G., HANSEN, J., Database Management And Design (2nd Edition), Prentice Hall, 1996
- [Ra02S] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, Slides for the 3rd Edition,
<http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Si19] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts (7th Edition), McGraw-Hill, 2019
- [Si19S] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts, Slides for the 7th Edition, <http://codex.cs.yale.edu/avi/db-book/>
- [Ul11] ULLMAN, J., WIDOM, J., A First Course in Database Systems,
<http://infolab.stanford.edu/~ullman/fcdb.html>