

Databases

Lectures 9-10*

The Physical Structure of Databases. Indexes

*30.11.2023 – Public holiday

Files of Records

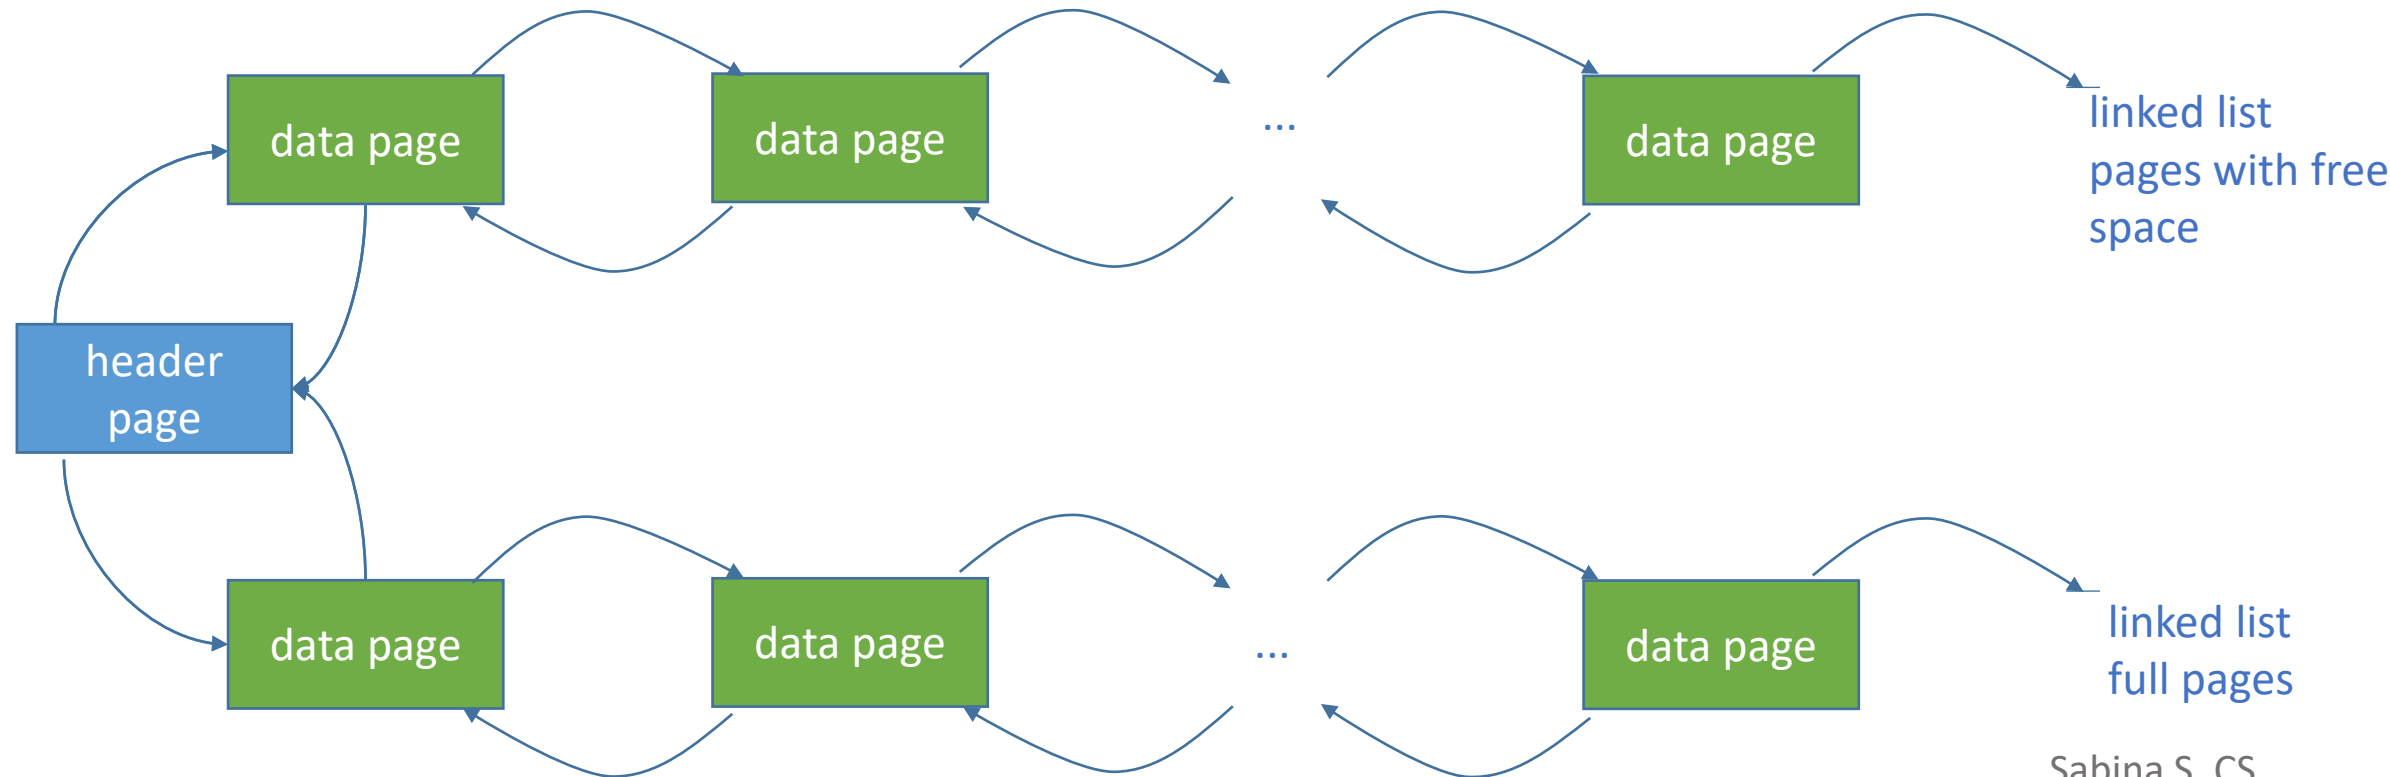
- higher level layers in the DBMS treat pages as collections of records
- file of records
 - collection of records; one or more pages
- different ways to organize a file's collection of pages
- every record has an identifier: the rid
- given the rid of a record, one can identify the page that contains the record

Heap Files

- the simplest file structure
- records are not ordered
- supported operations
 - create file
 - destroy file
 - insert a record
 - need to monitor pages with free space
 - retrieve a record given its rid
 - delete a record given its rid
 - scan all records
 - need to keep track of all the pages in the file
- appropriate when the expected pattern of use includes scans to obtain all the records

Heap Files - Linked List

- doubly linked list of pages
- DBMS stores the address of the first page (*header page*) of each file (a table holding pairs of the form $\langle \text{heap_file_name}, \text{page1_address} \rangle$)
- 2 lists – pages with free space and full pages




Heap Files - Linked List

- drawback
 - variable-length records => most of the pages will be in the list of pages with free space
 - when adding a record, multiple pages have to be checked until one is found that has enough free space

Heap Files - Directory of Pages

- DBMS stores the location of the header page for each heap file
- directory - collection of pages (e.g., linked list)
- directory entry - identifies a page in the file
- directory entry size - much smaller than the size of a page
- directory size - much smaller than the size of the file
- free space management
 - 1 bit / directory entry - corresponding page has / doesn't have free space
 - count / entry - available space on the corresponding page => efficient search of pages with enough free space when adding a variable-length record




Other File Organizations

- sorted files
 - suitable when data must be sorted, when doing range selections
- hashed files
 - files that are hashed on some fields (records are stored according to a hash function); good for equality selections

Record Formats

- fixed-length records



- each field has a fixed length
- fixed number of fields
- fields - stored consecutively
- computing a field's address
 - record address, length of preceding fields (from the system catalog)

Record Formats

- variable-length records
 - variable-length fields


v1

C ₁	\$	C ₂	\$	C ₃	\$
----------------	----	----------------	----	----------------	----

- fields
 - stored consecutively, separated by delimiters
- finding a field
 - a record scan

Record Formats

- variable-length records



- reserve space at the beginning of the record
 - array of fields offsets, offset to the end of the record
- array overhead, but direct access to every field

Page Formats


- page
 - collection of slots
 - 1 record / slot
- identifying a record
 - record id (rid): <page id, slot number>
- how to arrange records on pages
- how to manage slots

Page Formats

- fixed-length records
 - records have the same size
 - uniform, consecutive slots
 - adding a record
 - finding an available slot
- problems
 - keeping track of available slots
 - locating records


Page Formats

- fixed-length records - v1
 - n – number of records on the page
 - records are stored in the first n slots
 - locating record i - compute corresponding offset
 - deleting a record - the last record on the page is moved into the empty slot
 - empty slots - at the end of the page
- problems when a moved record has external references
 - the record's slot number would change, but the rid contains the slot number!



Page Formats

- fixed-length records - v2
- array of bits to monitor available slots
- 1 bit / slot
- deleting a record - turning off the corresponding bit




Page Formats

- variable-length records
 - adding a record
 - finding an empty slot of the right size
 - deleting a record
 - contiguous free space
 - a directory of slots / page
 - a pair <record offset , record length> / slot
 - a pointer to the beginning of the free space area on the page
 - moving a record on the page
 - only the record's offset changes
 - its slot number remains unmodified
 - can also be used for fixed-length records (when records need to be kept sorted)

Page Formats

- variable-length records



Indexes

- motivating example
 - file of students records sorted by name
 - good file organization
 - retrieve students in alphabetical order
 - not a good file organization
 - retrieve students whose age is in a given range
 - retrieve students who live in Timișoara
- index
 - auxiliary data structure that speeds up operations which can't be efficiently carried out given the file's organization
 - enables the retrieval of the rids of records that meet a selection condition (e.g., the rids of records describing students who live in Timișoara)

Indexes

- *search key*
 - set of one or more attributes of the indexed file (different from the *key* that identifies records)
- an index speeds up queries with equality / range selection conditions on the search key
- *entries*
 - records in the index (e.g., <search key, rid>)
 - enable the retrieval of records with a given search key value

Indexes

- example
 - files with students records
 - index built on attribute *city*
 - entries: $\langle \text{city}, \text{rid} \rangle$, where rid identifies a student record
 - such an index would speed up queries about students living in a given city:
 - find entries in the index with $\text{city} = \text{'Timișoara'}$
 - follow rids from obtained entries to retrieve records describing students who live in Timișoara

Indexes

- an index can improve the efficiency of certain types of queries, not of all queries (analogy - when searching for a book at the library, index cards sorted on author name cannot be used to efficiently locate a book given its title)
- organization techniques (access methods) - examples
 - B+ trees
 - hash-based structures
- changing the data in the file => update the indexes associated with the file (e.g., inserting records, updating search key columns, updating columns that are not part of the key, but are included in the index)
- index size
 - as small as possible, as indexes are brought into main memory for searches

Indexes - Data Entries


- problems
 - what does a data entry contain?
 - how are the entries of an index organized?
- let k^* be a data entry in an index; the data entry:
 - alternative 1
 - is an actual data record with search key value = k
 - alternative 2
 - is a pair $\langle k, \text{rid} \rangle$ (rid – id of a data record with search key value = k)
 - alternative 3
 - is a pair $\langle k, \text{rid_list} \rangle$ (rid_list – list of ids of data records with search key value = k)

Indexes - Data Entries

- a1
 - the file of data records needn't be stored in addition to the index
 - the index is seen as a special file organization
 - at most 1 index / collection of records should use alternative a1 (to avoid redundancy)
- a2, a3
 - data entries point to corresponding data records
 - in general, the size of an entry is much smaller than the size of a data record
 - a3 is more compact than a2, but can contain variable-length records
 - can be used by several indexes on a collection of records
 - independent of the file organization

Indexes - Data Entries

Ex. 1



Indexes - Data Entries

Ex. 1


- file with Employee records hashed on *age*
 - record <Ionescu, 60, F, 200>:
 - apply hash function to *age*: convert 60 to its binary representation, take the 2 least significant bits as the bucket identifier for the record
- index file that uses alternative 1 (data entries are the actual data records), search key *age*
- index that uses alternative 2 (data entries have the form <search key, rid>), search key *bonus*
- both indexes use hashing to locate data entries

Indexes - Data Entries

Ex. 2

	Name	Score	Age
rid ₁	Popescu	3	44
rid ₂	Ionescu	95	80
rid ₃	Vladescu	3	45
...		...	
rid _p	Xulescu	90	14

search key




data entry

...	
3	rid ₁ , rid ₃ ...
...	
90	rid _p ...
...	
95	rid ₂ ...
...	

index file

Clustered / Unclustered Indexes

- clustered index: the order of the data records is close to / the same as the order of the data entries
- unclustered index: index that is not clustered



References

- [Ta13] ȚÂMBULEA, L., Curs Baze de date, Facultatea de Matematică și Informatică, UBB, 2013-2014
- [Ra02] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (3rd Edition), McGraw-Hill, 2002
- [Ga09] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book (2nd Edition), Pearson Education, 2009
- [Ra02S] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, Slides for the 3rd Edition,
<http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Si19] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts (7th Edition), McGraw-Hill, 2019
- [Si19S] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts, Slides for the 7th Edition, <http://codex.cs.yale.edu/avi/db-book/>